

## USENIX Security 21

# ARCUS: Symbolic Root Cause Analysis of Exploits in Production Systems

论文分析 by 蒋隽

2022.07.24

# Contents

- 背景与矛盾
- 如何解决----Analyzing Root Cause Using Symbex
  - ARCUS Pipeline
  - Design
- Example: CVE-2018-12327
- 实验部分
  - 评估标准
  - 实验结果
- 结论与展望

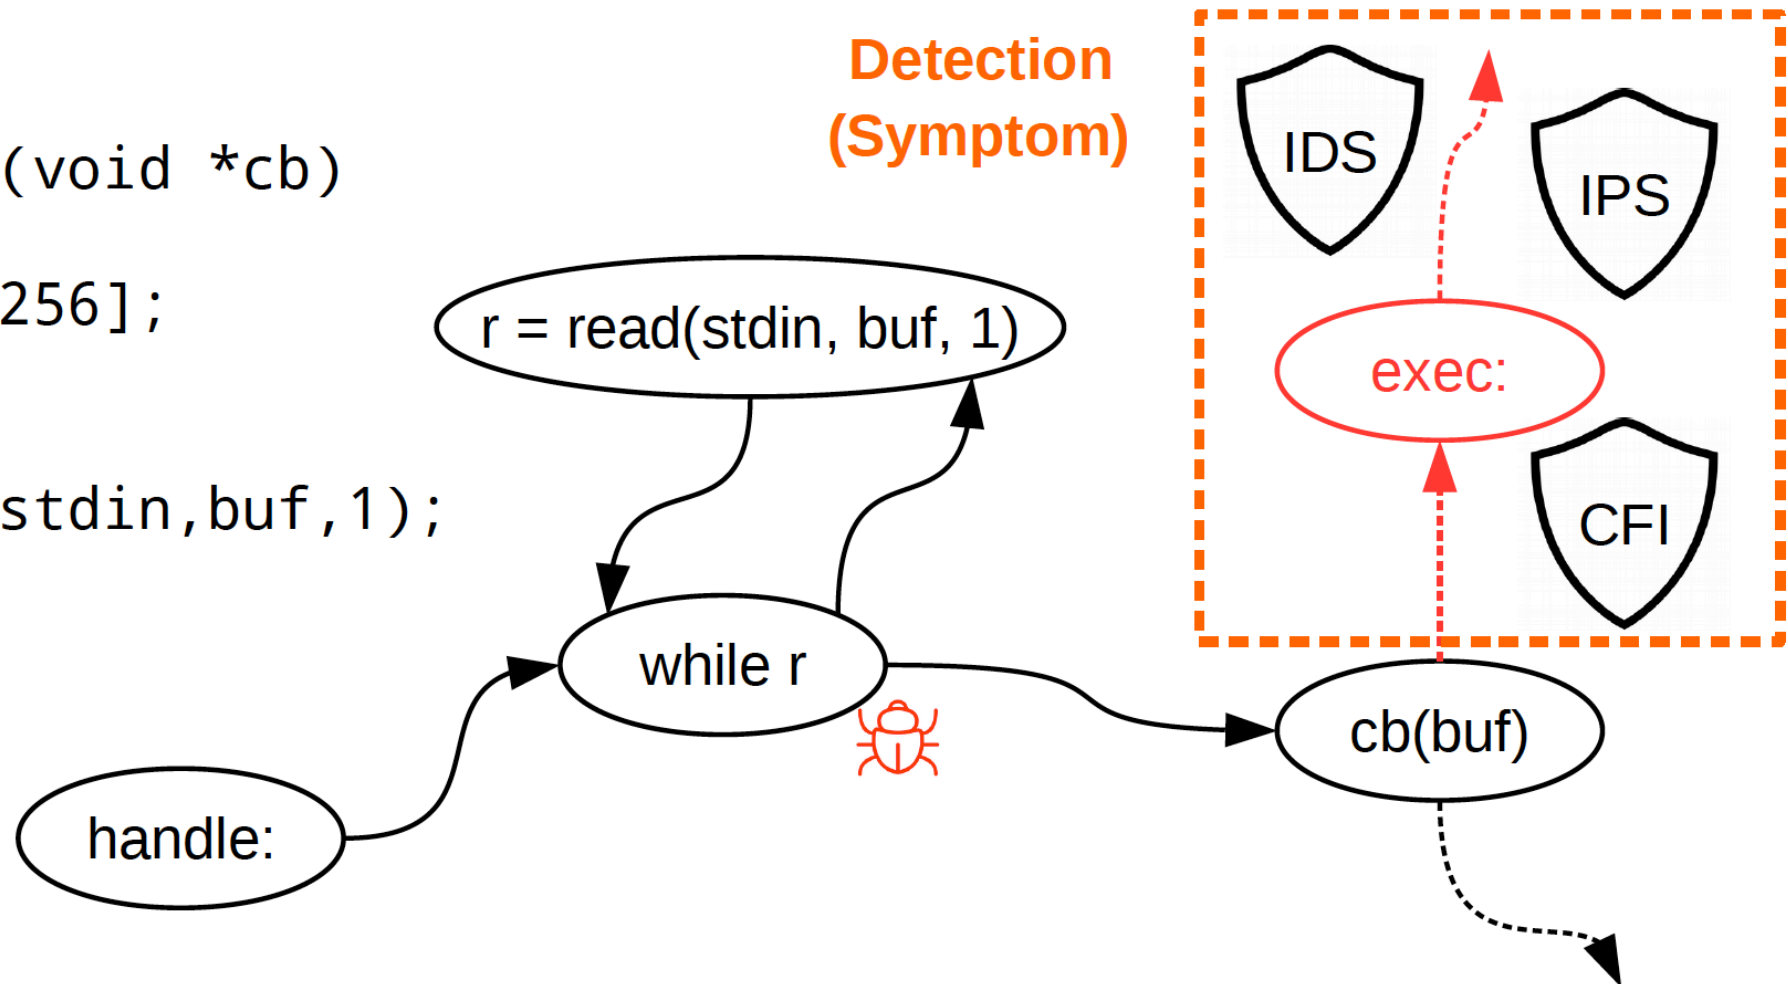
# Background

- 漏洞行为比漏洞形成的根本原因更容易被检测到
- 开发人员难以定位检测到安全风险的漏洞位置
- 发布补丁仅考虑表面修补而不是根本症结


**ARCUS** — — an automated framework for localizing the root cause of vulnerabilities in executions flagged by end-host runtime monitors.

# A Toy Case

```
void handle(void *cb)
{
    char buf[256];
    int r=1;
    while r
        r=read(stdin,buf,1);
    cb(buf);
}
```



# Why Symptoms?



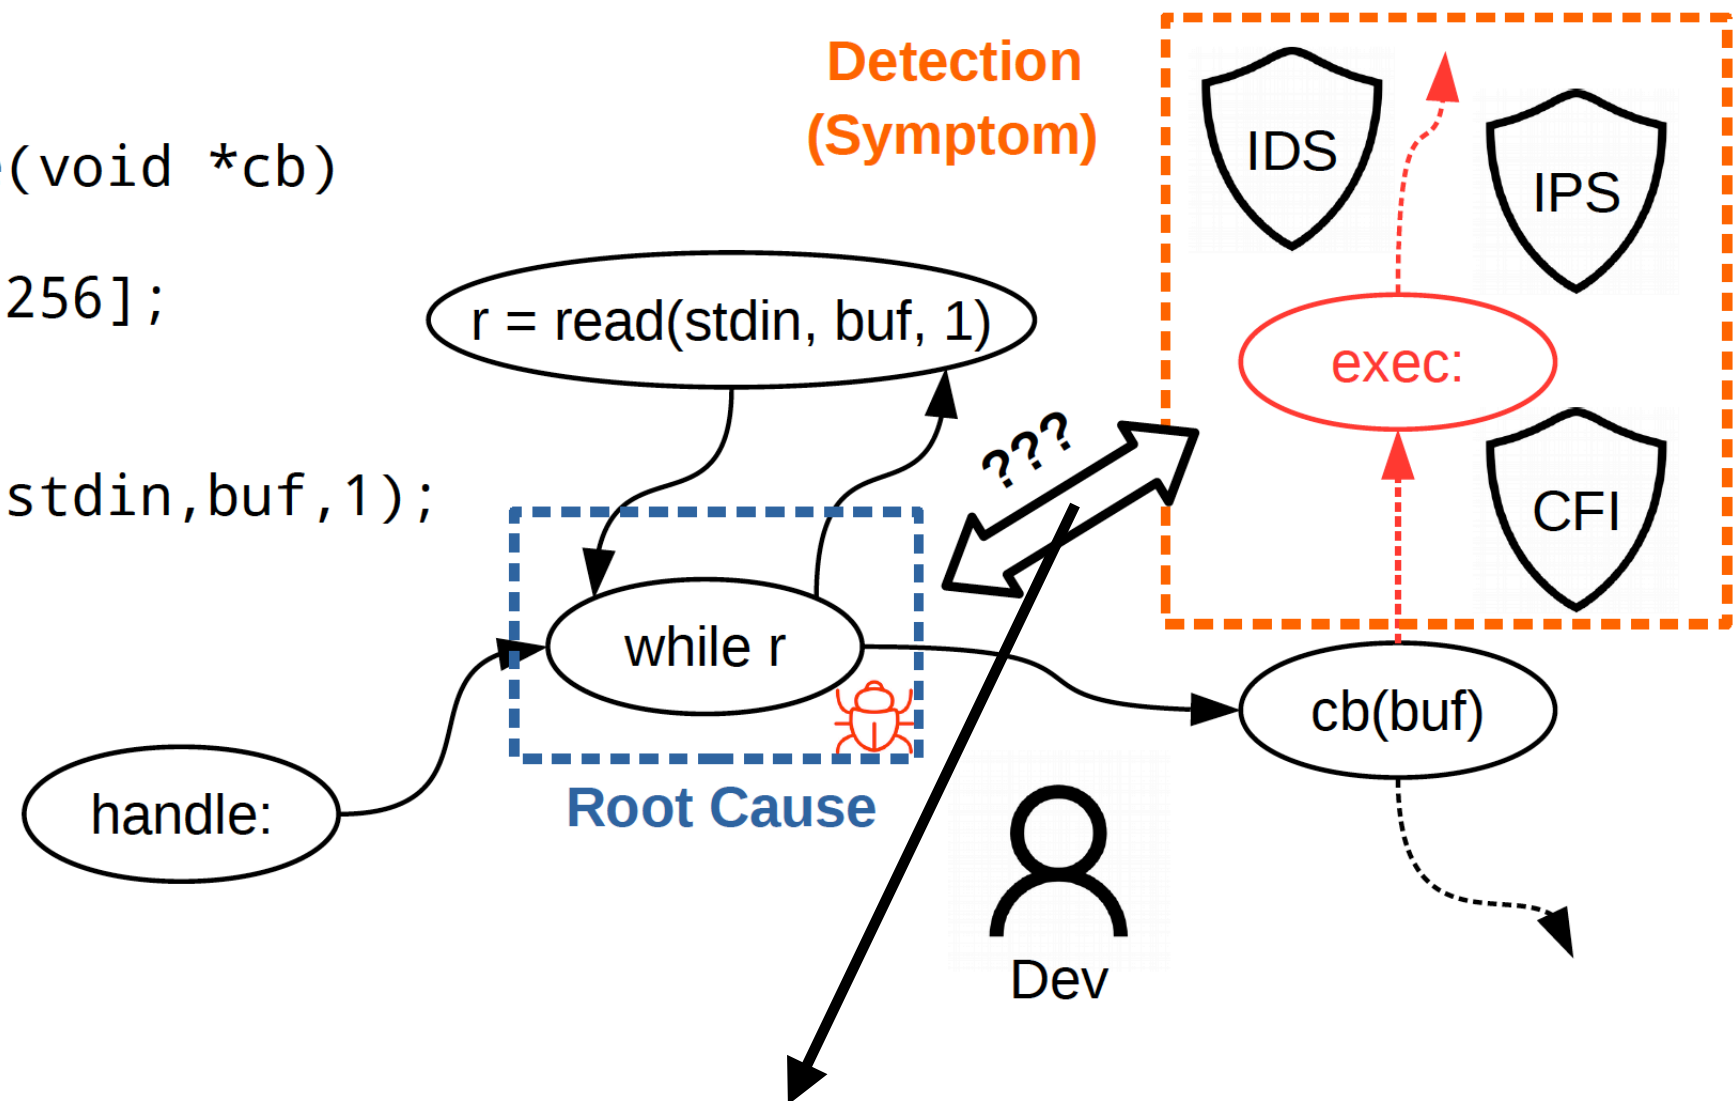
- Easiest to detect
  - Manifestation of behavior
- But how do we fix it?
  - Input filters
  - Function hardening
- Brittle, expensive

# Ideal Fix: Developer Patch

```
void handle(void *cb)
{
    char buf[256];
    int r=1;
    while r
        r=read(stdin,buf,1);
    cb(buf);
}
```

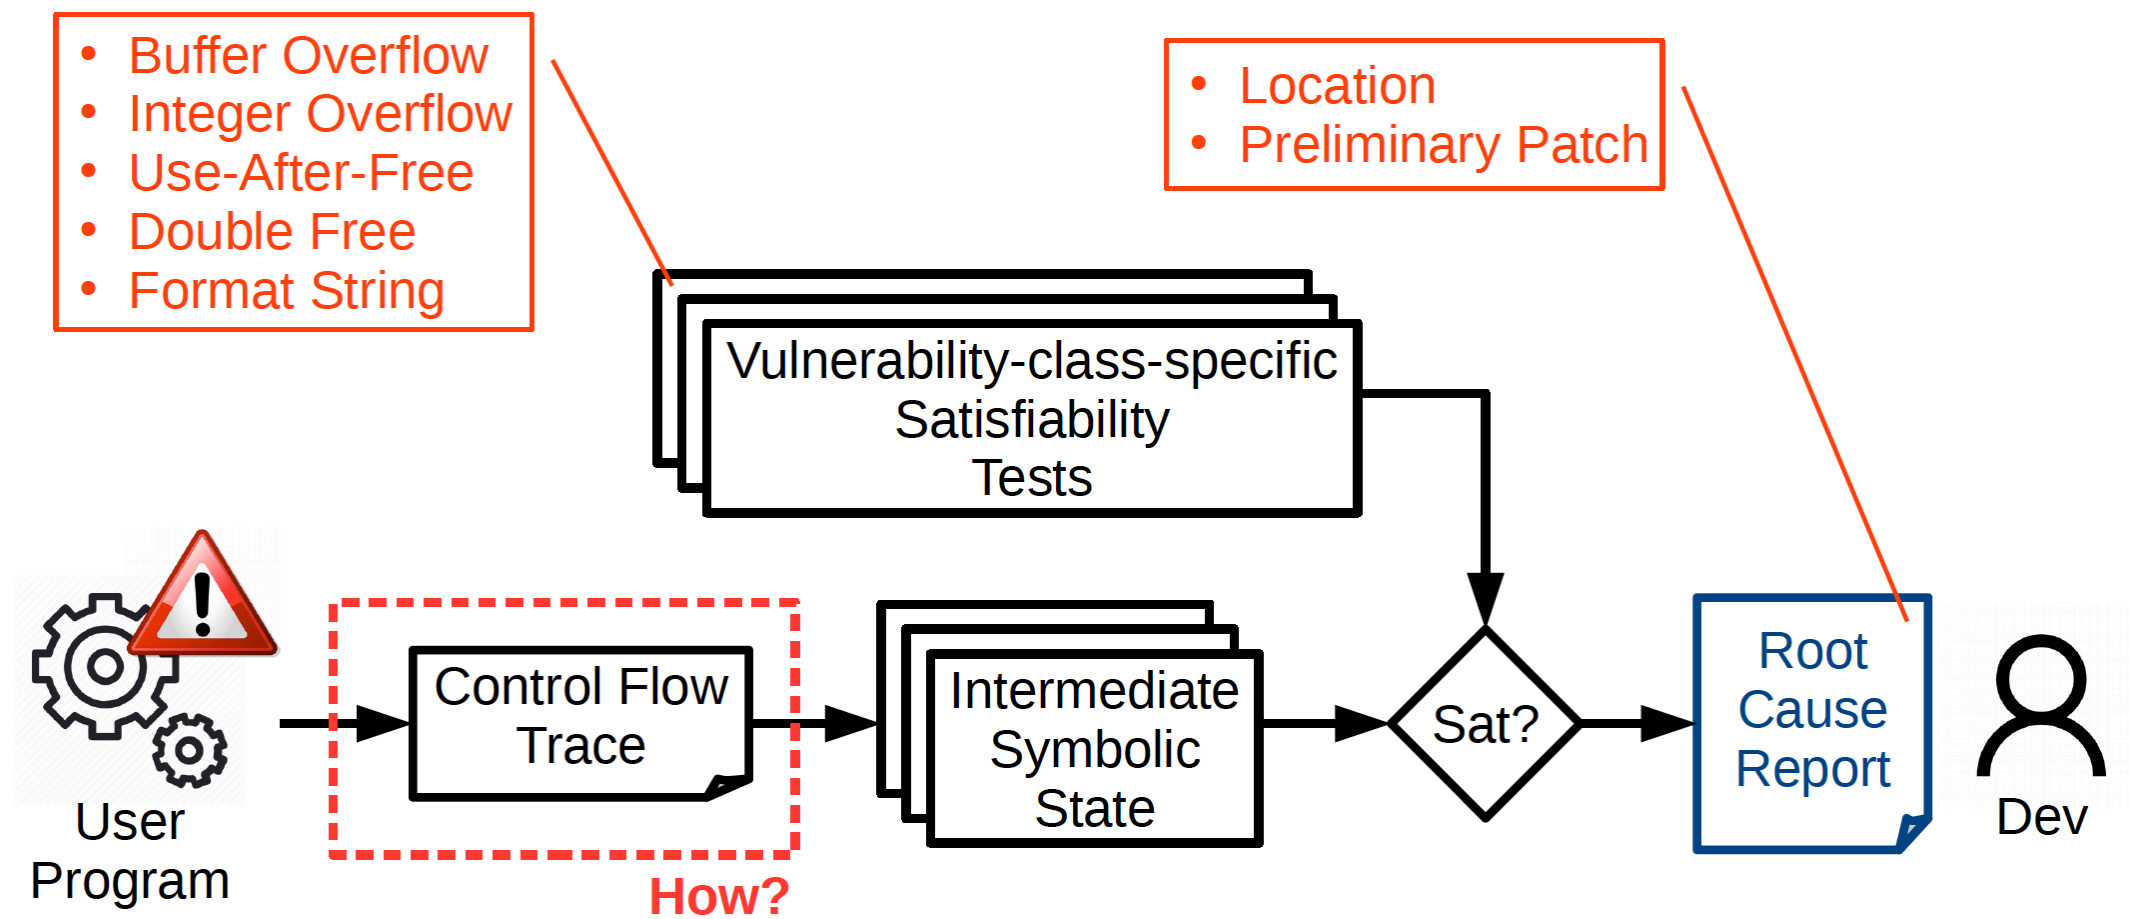


User  
Program



Real-World Cases Are Harder: Average Distance: 11,722 basic blocks

# ARCUS Pipeline



# Example: CVE-2018-12327

```
$ ./ntpq -4 [python -c 'print "A" * 300']
```

Name or service not known

```
*** stack smashing detected ***: <unknown> terminated
```



# The State Explosion Problem

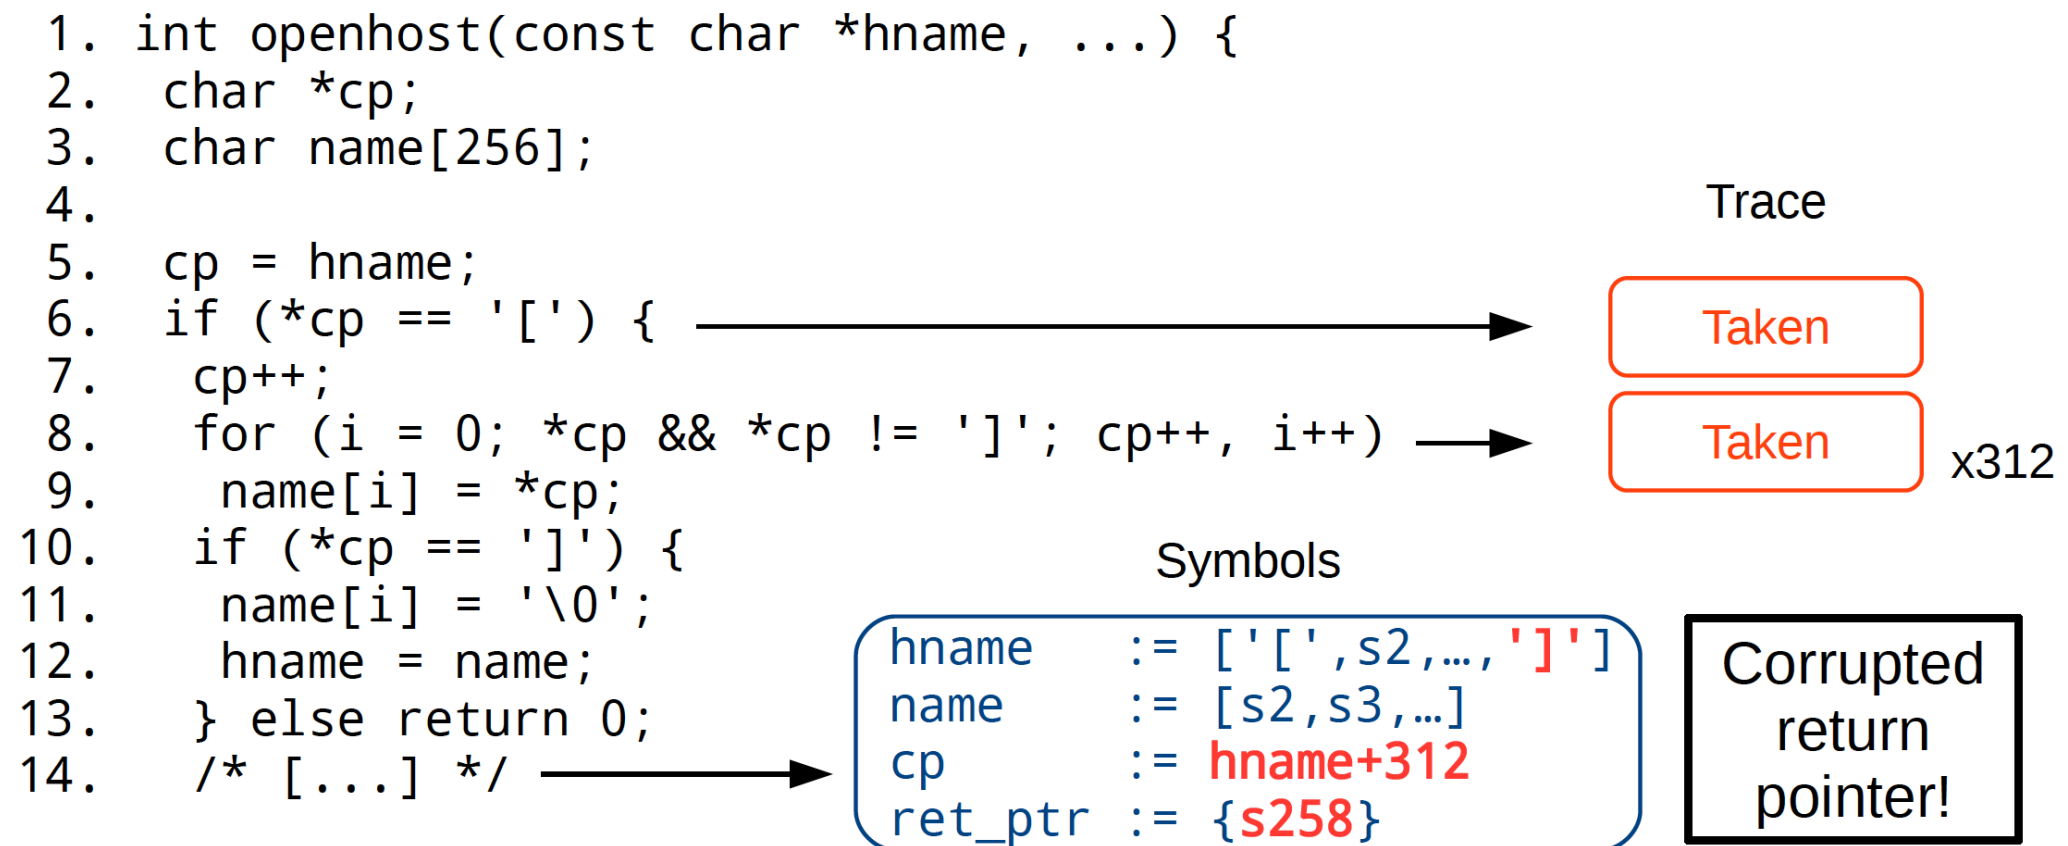
```
1. int openhost(const char *hname, ...) {  
2.   char *cp;   
3.   char name[256];  
4.  
5.   cp = hname;  
6.   if (*cp == '[') {  
7.     cp++;  
8.     for (i = 0; *cp && *cp != ']'; cp++, i++)  
9.       name[i] = *cp;  
10.    if (*cp == ']') {  
11.      name[i] = '\0';  
12.      hname = name;  
13.    } else return 0;  
14.    /* [...] */
```

Symbols

|         |    |             |
|---------|----|-------------|
| hname   | := | [s1,s2,...] |
| name    | := | []          |
| cp      | := | {}          |
| ret_ptr | := | {c1}        |

How many times  
should Line 9 iterate?

# Solution: Control Flow Trace



# Localizing Root Cause

Intermediate Symbolic States

```
hname  := ['[', s2, ...]  
name   := [s2]  
cp     := hname+1  
ret_ptr := {c1}
```


·  
·  
·



What if we *didn't*  
corrupt the pointer?

```
hname  := ['[', s2, ...]  
name   := [s2, s3, ...]  
cp     := hname+257  
ret_ptr := {c1}
```

exit loop



```
hname  := ['[', s2, ..., '']  
name   := [s2, s3, ...]  
cp     := hname+312  
ret_ptr := {s258}
```

```
hname  := ['[', s2, ..., '']  
name   := [s2, s3, ...]  
cp     := hname+257  
ret_ptr := {c1}
```

# What's Different?

```
hname  := ['[', s2, ..., ']']  
name   := [s2, s3, ...]  
cp     := hname+312  
ret_ptr := {s258}
```

$\Delta$

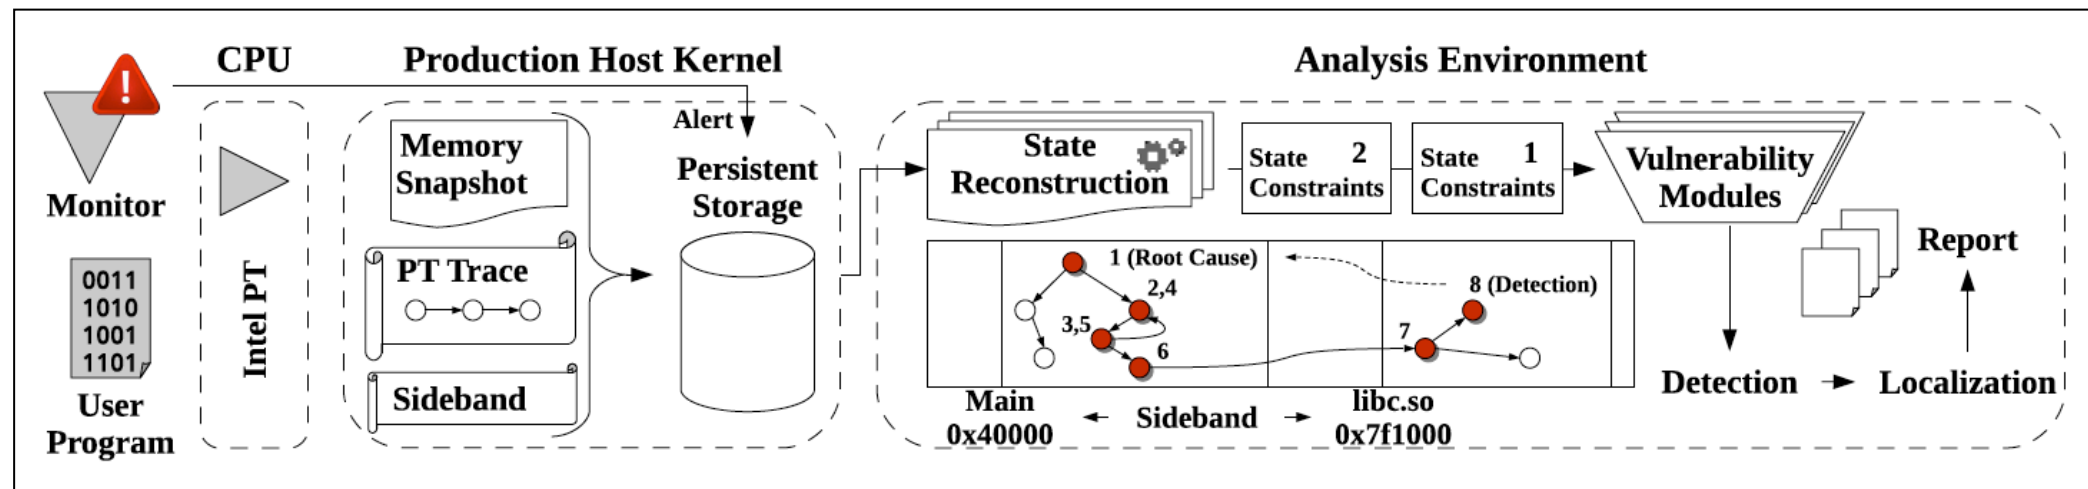
```
hname  := ['[', s2, ..., ']']  
name   := [s2, s3, ...]  
cp     := hname+257  
ret_ptr := {c1}
```

hname[257] != ']'  $\longleftrightarrow$  contradiction  $\longrightarrow$  hname[257] == ']'

Preliminary Patch:

```
for (i=0; *cp && *cp != ']'  
    && index(']', hname) <= 257; cp++, i++)
```

# ARCUS architecture



- 用户程序在终端主机中执行，而 ARCUS 内核模块使用 Intel PT 对其进行快照和跟踪。
- 当运行时监视器标记违规或异常时，数据将被发送到分析环境，在该环境中重建符号状态，模块在该环境中检测、定位和报告漏洞。

# Symbolic Execution Along Traced Paths

- 一旦监视器发出警报，ARCUS 将构建来自内核模块发送的数据的符号程序状态
- 使用符号分析，由于分析探索了不同的路径在程序中，它对符号数值施加约束，监测他们的数值
- 通过符号分析跟踪可以达到程序状态的可能数据值。
- 使用符号分析不是为了静态探索所有可能的路径，而是考虑所有可能的数据都流过一个特定的路径
- 符号化所有可以被攻击者控制的输入数据（命令行参数、环境变量、文件、套接字和其他标准 I/O），只为得到被追踪的路径，这避免了状态爆炸

# Analysis Modules — — Stack & Heap Overflow

*What if this path were to be taken by the program?*

- 识别变成符号的代码指针
- 确定编写它的基本块
- 找到控制写块执行的基本块
- 测试这些块上的附加约束是否可能使程序偏离错误行为

# Analysis Modules — — Stack & Heap Overflow

**Input:** VEX IR statements  $S$  starting from last executed.  
Tmp  $n$  to taint initially.

**Result:** Addresses  $A$  and registers  $R$  used to calculate  $n$ .

$A \leftarrow \emptyset$

$R \leftarrow \emptyset$

$T \leftarrow \{n\}$

**foreach**  $s$  in  $S$  **do**

**if**  $\text{Type}(s)$  is *Put* and  $\text{Type}(s.data)$  is *RdTmp* **then**

**if**  $s.data.tmp \in T$  **then**

$R \leftarrow R \cup \{s.register\}$

**end**

**end**

**if**  $\text{Type}(s)$  is *WrTmp* and  $s.tmp \in T$  **then**

**foreach**  $a$  in  $s.data.args$  **do**

**if**  $\text{Type}(a)$  is *Get* **then**

$R \leftarrow R \cup \{a.register\}$

**end**

**if**  $\text{Type}(a)$  is *RdTmp* **then**

$T \leftarrow T \cup \{a.tmp\}$

**end**

**if**  $\text{Type}(a)$  is *Load* **then**

$A \leftarrow A \cup \text{EvalTmp}(a.address)$

**end**

**end**

**end**

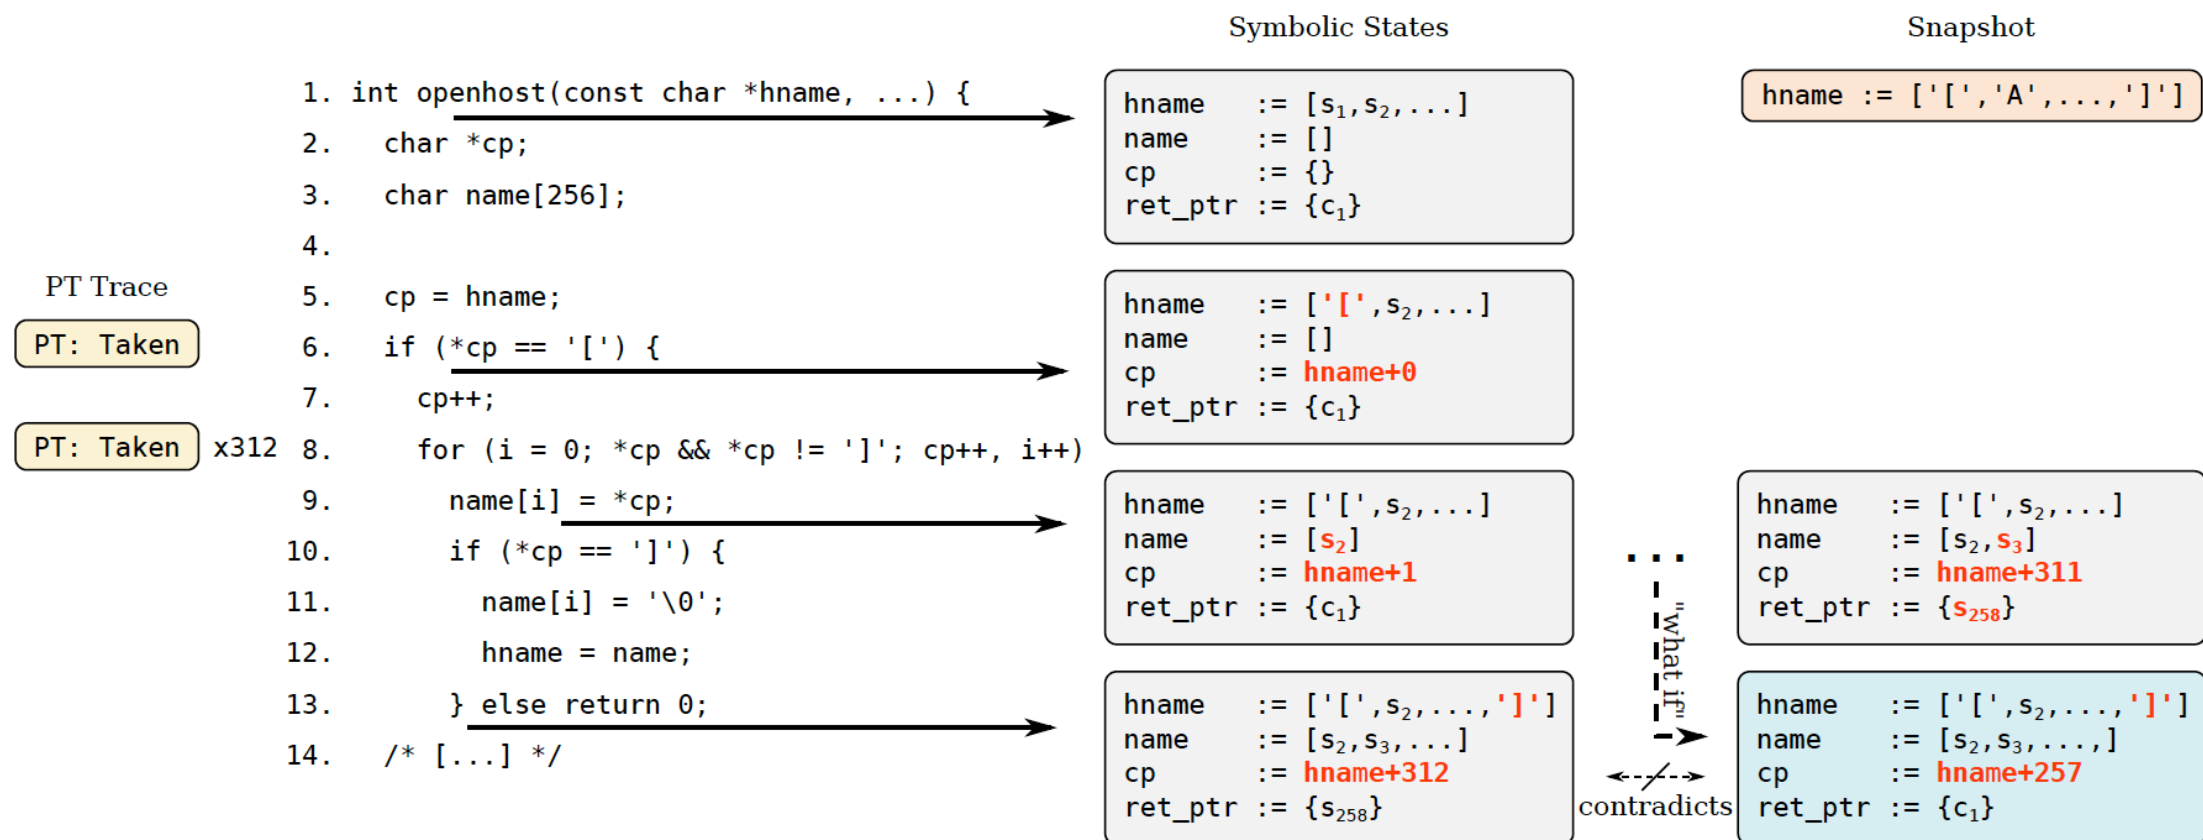
**end**

一旦被识别，该模块将通过先前重构的状态向后迭代，以找到包含在识别地址处的数据发生变化的状态。我们将此称为blame 状态。下一步是识别控制它的基本块(guardian)。该模块对重建状态使用前向分析来生成控制依赖图 (CDG) 并找到它们。(如左图)

如果有blame状态的guardian，则根据最短路径选择最接近的状态，并在先状态执行此操作重新访问代码以查看是否存在另一个分支，其约束与blame状态相矛盾。如果发现矛盾的约束，ARCUS 建议在guardian处强制执行。否则，只报告blame状态，因为需要一个全新的guardian。



# Revisiting CVE-2018-12327 in more detail.

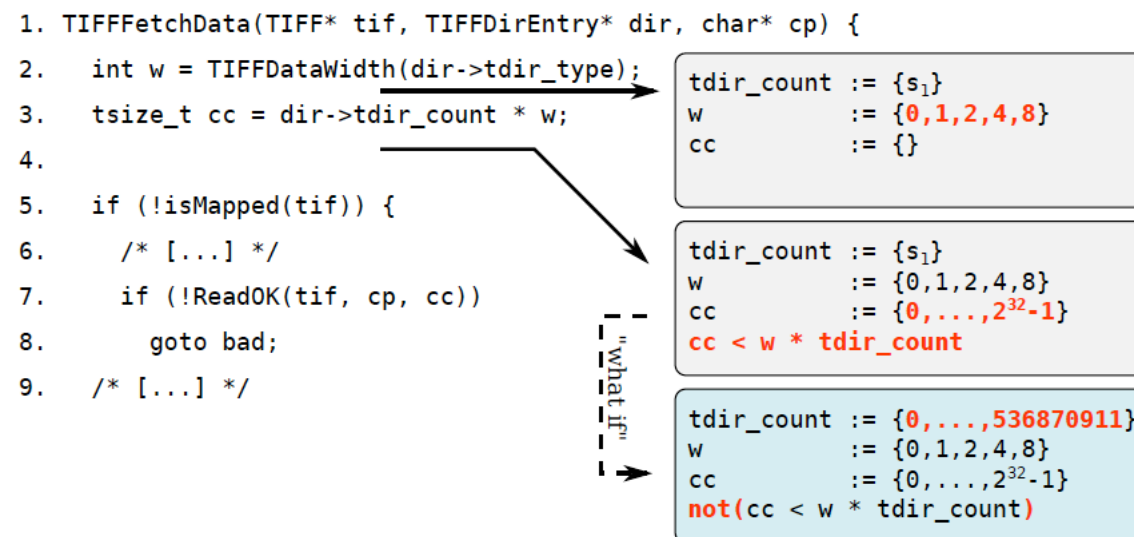


# Analysis Modules — Integer Overflow & Underflow

*What if the prior constraint was not satisfiable?*

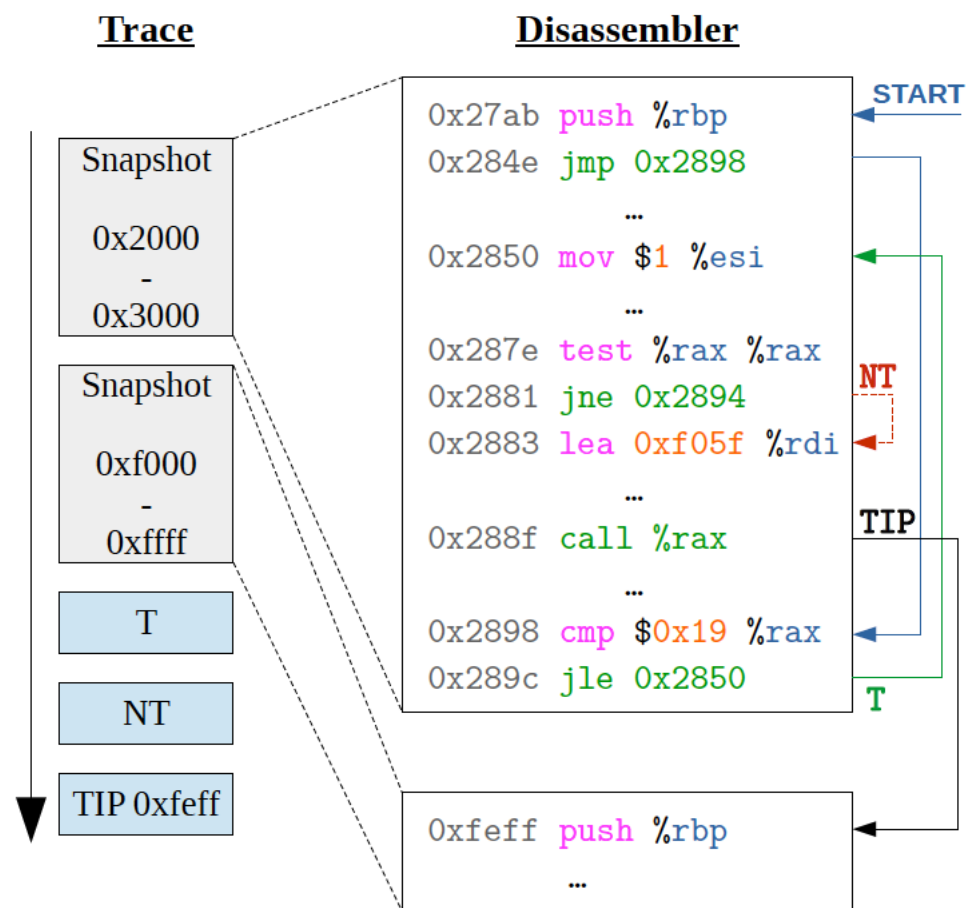
## TWO CHALLENGES

- 在没有类型信息的情况下推断寄存器和内存值的符号
- 避免由于开发人员和编译器故意溢出而导致的误报



在这种情况下，攻击者可以制作 TIFF 镜像以溢出保存 `cc` 的寄存器并将其传递给第 7 行的 `ReadOK`。由于 `cc` 是两个无符号值的乘积，因此 `cc < w * tdir_count` 在实际应用中是不可能的，但在第 4 行模块发现它是可满足的，表明 `cc` 可以溢出。当 `cc` 然后传递给 `ReadOK` 时，模块会标记错误。

# Capturing the Executed Path



- 采用内核模块管理PT
- 英特尔 PT假设审计员知道已审核的程序，我们的内核模块将其添加到Trace作为快照（灰色数据包）
- 内核模块还捕获和插入在 PT 数据之间动态生成代码页（蓝色地址）
- 对于条件分支，单个采用-未采用位被记录下来。对于间接控制流传输（返回、间接调用和间接跳转）和异步事件（例如，中断、异常），记录目的地
- 英特尔 PT 绕过所有缓存和记忆翻译，最大限度地减少其对跟踪的程序。当为跟踪分配的缓冲区被填满时，CPU 引发一个不可屏蔽中断 (NMI)，内核模块立即处理，因此不会丢失任何数据

# Evaluation — — Several Questions

- Is ARCUS **accurate** at detecting bugs within our covered classes?
- Can ARCUS locate and analyse **real-world exploits**?
- Are ARCUS' root cause reports consistent with **real world advisories and patches**?
- Is ARCUS **feasible** to deploy in terms of runtime and storage overhead?

- Is ARCUS **accurate** at detecting bugs within our covered classes?

| Table 3: RIPE and Juliet Test Cases                |       |       |    |     |       |
|--|-------|-------|----|-----|-------|
| Overall Results (Detection by $\geq 1$ Strategies) |       |       |    |     |       |
| RIPE   | TP    | TN    | FP | FN  | Acc.  |
| BSS  | 170   | 170   | 0  | 0   | 100%  |
| Data   | 190   | 190   | 0  | 0   | 100%  |
| Heap   | 190   | 190   | 0  | 0   | 100%  |
| Stack  | 260   | 260   | 0  | 0   | 100%  |
| Juliet   | TP    | TN    | FP | FN  | Acc.  |
| CWE-134  | 1,200 | 2,600 | 0  | 0   | 100%  |
| CWE-415  | 818   | 2,212 | 0  | 0   | 100%  |
| CWE-416  | 393   | 1,222 | 0  | 0   | 100%  |
| By Locating Strategy (RIPE)                        |       |       |    |     |       |
| Symbolic IP  | TP    | TN    | FP | FN  | Acc.  |
| BSS  | 154   | 170   | 0  | 16  | 95.3% |
| Data   | 171   | 190   | 0  | 19  | 95.0% |
| Heap   | 154   | 190   | 0  | 36  | 90.5% |
| Stack  | 211   | 260   | 0  | 49  | 90.6% |
| Int Overflow                                       | TP    | TN    | FP | FN  | Acc.  |
| BSS  | 60    | 170   | 0  | 110 | 67.6% |
| Data   | 60    | 190   | 0  | 130 | 65.8% |
| Heap   | 60    | 190   | 0  | 130 | 65.8% |
| Stack  | 150   | 260   | 0  | 110 | 78.8% |
| By Locating Strategy (Juliet)                      |       |       |    |     |       |
| Symbolic Args.                                     | TP    | TN    | FP | FN  | Acc.  |
| CWE-134  | 1,200 | 2,600 | 0  | 0   | 100%  |
| Track Frees  | TP    | TN    | FP | FN  | Acc.  |
| CWE-415  | 818   | 2,212 | 0  | 0   | 100%  |
| R/W Freed Addrs.                                   | TP    | TN    | FP | FN  | Acc.  |
| CWE-416  | 393   | 1,222 | 0  | 0   | 100%  |

- ARCUS 正确分析了所有套件中的所有测试用例，没有 FP 或 FN。也就是说，每个 TP 至少被 1 个模块检测到，而 TN 没有被检测到。我们手动验证 TP 案例的根本原因报告是否正确识别了有缺陷的功能，并且建议可以防止内存损坏
- ARCUS 在 RIPE 案例上非常准确，因为有多种机会检测溢出。例如，损坏返回指针的整数溢出可以在寄存器回绕时由整数溢出模块检测，或者在指针被覆盖时由堆栈溢出模块检测
- 对于 Juliet 案例测试的模块，它们的能力不重叠并且产生相同的数字和整个表格一样。对于与 RIPE 相关的策略，我们发现符号 IP 检测的平均准确率为 92.9%，而整数溢出检测的准确率为 69.5%

- Can ARCUS locate and analyse **real-world exploits**?

| CVE / EDB        | Type    | Program             | # BBs            | Size (MB)   | ΔRoot Cause   | ΔAlert        | Located | Has Patch | Match            |
|------------------|---------|---------------------|------------------|-------------|---------------|---------------|---------|-----------|------------------|
| CVE-2004-0597    | Heap    | GIMP (libpng)       | 41,625,163       | 56.0        | 247           | 1             | Yes     | [61]      | Yes <sup>†</sup> |
| CVE-2004-1279    | Heap    | jpegtoavi           | 67,772           | 0.65        | 26,216        | 1             | Yes     | No        | -                |
| CVE-2004-1288    | Heap    | o3read              | 74,723           | 0.65        | 33,211        | 1             | Yes     | [62]      | Yes              |
| CVE-2009-2629    | Heap    | nginx               | 300,071          | 1.10        | 28            | 33,824        | Yes     | [63]      | Yes              |
| CVE-2009-3896    | Heap    | nginx               | 283,157          | 1.10        | 59            | 16,821        | Yes     | [64]      | Yes              |
| CVE-2017-9167    | Heap    | autotrace           | 75,404           | 1.01        | 1,828         | 2             | Yes     | No        | -                |
| CVE-2018-12326   | Heap    | Redis               | 291,275          | 1.20        | 8             | 234           | Yes     | [65]      | Yes              |
| EDB-15705        | Heap    | ftp                 | 260,986          | 0.85        | 19,322        | 2             | Yes     | No        | -                |
| CVE-2004-1257    | Stack   | abc2mtex            | 53,490           | 0.67        | 6,319         | 1             | Yes     | No        | -                |
| CVE-2009-5018    | Stack   | gif2png             | 90,738           | 1.09        | 1,848         | 1             | Yes     | [66]      | Yes              |
| CVE-2017-7938    | Stack   | dmitry              | 100,186          | 0.71        | 4,051         | 14,402        | Yes     | No        | -                |
| CVE-2018-12327   | Stack   | ntpq                | 374,830          | 1.85        | 122,740       | 77,990        | Yes     | [67]      | Yes              |
| CVE-2018-18957   | Stack   | GOOSE (libiec61850) | 65,198           | 0.71        | 94            | 30            | Yes     | [68]      | Yes              |
| CVE-2019-14267   | Stack   | pdfresurrect        | 128,427          | 0.66        | 83,123        | 1             | Yes     | [69]      | Yes              |
| * EDB-47254      | Stack   | abc2mtex            | 53,490           | 0.67        | 6,566         | -             | Yes     | No        | -                |
| EDB-46807        | Stack   | MiniFtp             | 60,849           | 0.69        | 335           | 107           | Yes     | No        | -                |
| CVE-2006-2025    | Integer | GIMP (libtiff)      | 78,419,067       | 55.0        | 3             | 8             | Yes     | [70]      | Yes              |
| CVE-2007-2645    | Integer | exif (libexif)      | 67,697           | 0.97        | 1             | 7             | Yes     | [71]      | Yes              |
| CVE-2013-2028    | Integer | nginx               | 809,977          | 2.00        | 1             | 25,268        | Yes     | [72]      | Yes              |
| CVE-2017-7529    | Integer | nginx               | 1,049,494        | 1.10        | 2             | 780,404       | Yes     | [73]      | Yes              |
| CVE-2017-9186    | Integer | autotrace           | 75,142           | 1.00        | 1             | 1             | Yes     | No        | -                |
| CVE-2017-9196    | Integer | autotrace           | 74,695           | 1.03        | 1             | 203           | Yes     | No        | -                |
| * CVE-2019-19004 | Integer | autotrace           | 132,302          | 1.02        | 1             | -             | Yes     | No        | -                |
| CVE-2017-11403   | UAF     | GraphicsMagick      | 2,316,152        | 4.61        | 38            | 1             | Yes     | [74]      | Yes              |
| CVE-2017-14103   | UAF     | GraphicsMagick      | 2,316,133        | 4.61        | 38            | 1             | Yes     | [74]      | Yes              |
| CVE-2017-9182    | UAF     | autotrace           | 132,302          | 1.02        | 296           | 58,058        | Yes     | No        | -                |
| * CVE-2019-17582 | UAF     | PHP (libzip)        | 5,980,255        | 6.40        | 49            | -             | Yes     | [75]      | Yes              |
| CVE-2017-12858   | DF      | PHP (libzip)        | 5,980,255        | 6.40        | 51            | 719           | Yes     | [75]      | Yes              |
| * CVE-2019-19005 | DF      | autotrace           | 132,302          | 1.02        | 57,859        | -             | Yes     | No        | -                |
| CVE-2005-0105    | FS      | typespeed           | 127,209          | 0.74        | 1             | 1             | Yes     | [76]      | Yes              |
| CVE-2012-0809    | FS      | sudo                | 108,442          | 0.69        | 1             | 1             | Yes     | [77]      | Yes              |
| <b>Average:</b>  |         |                     | <b>4,568,619</b> | <b>5.07</b> | <b>11,722</b> | <b>36,804</b> |         |           |                  |


\* New vulnerability discovered by ARCUS.

<sup>†</sup> Equivalent to applied patch.

- Are ARCUS' root cause reports consistent with **real world advisories and patches**?
- 通过手动将它们与公共漏洞进行比较评估漏洞利用报告的质量
- 在这些案例中，我们使用补丁来进一步验证 ARCUS 的质量通过手动确认他们识别出相同的代码



- Is ARCUS **feasible** to deploy in terms of runtime and storage overhead?



跟踪 SPEC CPU 基准的性能开销和存储大小。 平均开销为 7.21%，  
几何平均值为 3.81%。 平均跟踪大小为 110 MB，几何平均值为 38.2 MB。



# Results

- 我们证明我们的方法可以构建符号程序状态并分析几类严重和普遍的软件漏洞
- 我们对 27 个漏洞和 9,000 多个 Juliet 和 RIPE 测试用例的评估表明, ARCUS 可以自动识别所有测试漏洞的根本原因, 在此过程中发现 4 个新漏洞
- ARCUS 在 SPEC 2006 CPU 基准测试中产生了 7.21% 的性能开销, 并可扩展到从超过 810,000 行 C/C++ 代码编译的大型程序

***Thanks***

Code: <https://github.com/carter-yagemann/ARCUS>